

# A Hierarchical Approach for Generating Descriptive Image Paragraphs

Jonathan Krause   Justin Johnson   Ranjay Krishna   Li Fei-Fei  
Stanford University

{jkrause, jcjohns, ranjaykrishna, feifeili}@cs.stanford.edu

## Abstract

Recent progress on image captioning has made it possible to generate novel sentences describing images in natural language, but compressing an image into a single sentence can describe visual content in *only coarse detail*. While one new captioning approach, dense captioning, can potentially describe images in finer levels of detail by captioning many regions within an image, it in turn is unable to produce a *coherent* story for an image. In this paper we overcome these limitations by generating entire paragraphs for describing images, which can tell detailed, unified stories. We develop a model that decomposes both images and paragraphs into their constituent parts, detecting semantic regions in images and using a *hierarchical recurrent neural network* to reason about language. Linguistic analysis confirms the complexity of the paragraph generation task, and thorough experiments on a new dataset of image and paragraph pairs demonstrate the effectiveness of our approach.

## 1. Introduction

Vision is the primary sensory modality for human perception, and language is our most powerful tool for communicating with the world. Building systems that can simultaneously understand visual stimuli and describe them in natural language is therefore a core problem in both computer vision and artificial intelligence as a whole. With the advent of large datasets pairing images with natural language descriptions [20, 34, 10, 16] it has recently become possible to generate novel sentences describing images [4, 6, 12, 22, 30]. While the success of these methods is encouraging, they all share one key limitation: *detail*. By only describing images with a single high-level sentence, there is a fundamental upper-bound on the quantity and quality of information approaches can produce.

One recent alternative to sentence-level captioning is the task of *dense captioning* [11], which overcomes this limitation by detecting many regions of interest in an image and describing each with a short phrase. By extending the task of object detection to include natural language description,



### Sentences

- 1) A girl is eating donuts with a boy in a restaurant
- 2) A boy and girl sitting at a table with doughnuts.
- 3) Two kids sitting a coffee shop eating some frosted donuts
- 4) Two children sitting at a table eating donuts.
- 5) Two children eat doughnuts at a restaurant table.

### Paragraph

Two children are sitting at a table in a restaurant. The children are one little girl and one little boy. The little girl is eating a pink frosted donut with white icing lines on top of it. The girl has blonde hair and is wearing a green jacket with a black long sleeve shirt underneath. The little boy is wearing a black zip up jacket and is holding his finger to his lip but is not eating. A metal napkin dispenser is in between them at the table. The wall next to them is white brick. Two adults are on the other side of the short white brick wall. The room has white circular lights on the ceiling and a large window in the front of the restaurant. It is daylight outside.

Figure 1. Paragraphs are longer, more informative, and more linguistically complex than sentence-level captions. Here we show an image with its sentence-level captions from MS COCO [20] (top) and the paragraph used in this work (bottom).

dense captioning describes images in considerably more detail than standard image captioning. However, this comes at a cost: *descriptions generated for dense captioning are not coherent*, *i.e.* they do not form a cohesive whole describing the entire image.

In this paper we address the shortcomings of both traditional image captioning and the recently-proposed dense

image captioning by introducing the task of generating paragraphs that richly describe images (Fig. 1). Paragraph generation combines the strengths of these tasks but does not suffer from their weaknesses – like traditional captioning, paragraphs give a coherent natural language description for images, but like dense captioning, they can do so in fine-grained detail.

Generating paragraphs for images is challenging, requiring both fine-grained image understanding and long-term language reasoning. To overcome these challenges, we propose a model that decomposes images and paragraphs into their constituent parts: We break images into semantically meaningful pieces by detecting objects and other regions of interest, and we reason about language with a hierarchical recurrent neural network, decomposing paragraphs into their corresponding sentences. In addition, we also demonstrate for the first time the ability to transfer visual and linguistic knowledge from large-scale region captioning [16], which we show has the ability to improve paragraph generation.

To validate our method, we collected a dataset of image and paragraph pairs, which complements the whole-image and region-level annotations of MS COCO [20] and Visual Genome [16]. To validate the complexity of the paragraph generation task, we performed a linguistic analysis of our collected paragraphs, comparing them to sentence-level image captioning. We compare our approach with numerous baselines, showcasing the benefits of hierarchical modeling for generating descriptive paragraphs.

The rest of this paper is organized as follows: Sec. 2 overviews related work in image captioning and hierarchical RNNs, Sec. 3 introduces the paragraph generation task, describes our newly-collected dataset, and performs a simple linguistic analysis on it, Sec. 4 details our model for paragraph generation, Sec. 5 contains experiments, and Sec. 6 concludes with discussion.

## 2. Related Work

**Image Captioning** Building connections between visual and textual data has been a longstanding goal in computer vision. One line of work treats the problem as a ranking task, using images to retrieve relevant captions from a database and vice-versa [8, 10, 13]. Due to the compositional nature of language, it is unlikely that any database will contain all possible image captions; therefore another line of work focuses on generating captions directly. Early work uses handwritten templates to generate language [17] while more recent methods train recurrent neural network language models conditioned on image features [4, 6, 12, 22, 30, 33] and sample from them to generate text. Similar methods have also been applied to generate captions for videos [6, 32, 35].

A handful of approaches to image captioning reason not only about whole images but also image regions. Xu *et al.* [31] generate captions using a recurrent network with

attention, so that the model produces a distribution over image regions for each word. In contrast to their work, which uses a coarse grid as image regions, we use semantically meaningful regions of interest. Karpathy and Fei-Fei [12] use a ranking loss to align image regions with sentence fragments but do not do generation with the model. Johnson *et al.* [11] introduce the task of dense captioning, which detects and describes regions of interest, but these descriptions are independent and do not form a coherent whole.

There has also been some pioneering work on video captioning with multiple sentences [27]. While videos are a natural candidate for multi-sentence description generation, image captioning cannot leverage strong temporal dependencies, adding extra challenge.

**Hierarchical Recurrent Networks** In order to generate a paragraph description, a model must reason about long-term linguistic structures spanning multiple sentences. Due to vanishing gradients, recurrent neural networks trained with stochastic gradient descent often struggle to learn long-term dependencies. Alternative recurrent architectures such as long-short term memory (LSTM) [9] help alleviate this problem through a gating mechanism that improves gradient flow. Another solution is a *hierarchical* recurrent network, where the architecture is designed such that different parts of the model operate on different time scales.

Early work applied hierarchical recurrent networks to simple algorithmic problems [7]. The Clockwork RNN [15] uses a related technique for audio signal generation, spoken word classification, and handwriting recognition; a similar hierarchical architecture was also used in [2] for speech recognition. In these approaches, each recurrent unit is updated on a fixed schedule: some units are updated on every timestep, while other units might be updated every other or every fourth timestep. This type of hierarchy helps reduce the vanishing gradient problem, but the hierarchy of the model does not directly reflect the hierarchy of the output sequence.

More related to our work are hierarchical architectures that directly mirror the hierarchy of language. Li *et al.* [18] introduce a hierarchical autoencoder, and Lin *et al.* [19] use different recurrent units to model sentences and words. Most similar to our work is Yu *et al.* [35], who generate multi-sentence descriptions for cooking videos using a different hierarchical model. Due to the less constrained non-temporal setting in our work, our method has to learn in a much more generic fashion and has been made simpler as a result, relying more on learning the interplay between sentences. Additionally, our method reasons about semantic regions in images, which both enables the transfer of information from these regions and leads to more interpretability in generation.

	Sentences COCO [20]	Paragraphs Ours
Description Length	11.30	67.50
Sentence Length	11.30	11.91
Diversity	19.01	70.49
Nouns	33.45%	25.81%
Adjectives	27.23%	27.64%
Verbs	10.72%	15.21%
Pronouns	1.23%	2.45%

Table 1. Statistics of paragraph descriptions, compared with sentence-level captions used in prior work. Description and sentence lengths are represented by the number of tokens present, diversity is the inverse of the average CIDEr score between sentences of the same image, and part of speech distributions are aggregated from Penn Treebank [23] part of speech tags.

### 3. Paragraphs are Different

To what extent does describing images with paragraphs differ from sentence-level captioning? To answer this question, we collected a **novel dataset of paragraph annotations**, comprised of 19,551 MS COCO [20] and Visual Genome [16] images, where each image has been annotated with a paragraph description. Annotations were collected on Amazon Mechanical Turk, using U.S. workers with at least 5,000 accepted HITs and an acceptance rate of 98% or greater<sup>1</sup>, and were additionally subject to automatic and manual spot checks on quality. Fig. 1 demonstrates an example, comparing our collected paragraph with the five corresponding sentence-level captions from MS COCO. Though it is clear that the paragraph is longer and more descriptive than any one sentence, we note further that a single paragraph can be more detailed than *all five* sentence captions, even when combined. This occurs because of redundancy in sentence-level captions – while each caption might use slightly different words to describe the image, since **all sentence captions have the goal of describing the image as a whole**, they are fundamentally limited in terms of both diversity and their total information.

We quantify these observations along with various other statistics of language in Tab. 1. For example, we find that each paragraph is roughly six times as long as the average sentence caption, and the individual sentences in each paragraph are of comparable length as sentence-level captions. To examine the issue of sentence diversity, we compute the average CIDEr [29] similarity between COCO sentences for each image and between the individual sentences in each collected paragraph, defining the final diversity score as 100 minus the average CIDEr similarity. Viewed through this metric, the difference in diversity is striking – sentences

within paragraphs are substantially more diverse than sentence captions, with a **diversity score of 70.49 compared to only 19.01**. This quantifiable evidence demonstrates that sentences in paragraphs provide significantly more information about images.

Diving deeper, we performed a simple linguistic analysis on COCO sentences and our collected paragraphs, comprised of annotating each word with a part of speech tag from Penn Treebank via Stanford CoreNLP [21] and aggregating parts of speech into higher-level linguistic categories. A few common parts of speech are given in Tab. 1. As a proportion, paragraphs have somewhat more verbs and pronouns, a comparable frequency of adjectives, and somewhat fewer nouns. Given the nature of paragraphs, this makes sense – longer descriptions go beyond the presence of a few salient objects and include information about their properties and relationships. We also note but do not quantify that paragraphs exhibit higher frequencies of more complex linguistic phenomena, *e.g.* coreference occurring in Fig. 1, wherein sentences refer to either “two children”, “one little girl and one little boy”, “the girl”, or “the boy.” We believe that these types of long-range phenomena are a fundamental property of descriptive paragraphs with human-like language and cannot be adequately explored with sentence-level captions.

## 4. Method

**Overview** Our model takes an image as input, generating a natural-language paragraph describing it, and is designed to take advantage of the compositional structure of both images and paragraphs. Fig. 2 provides an overview. We first decompose the input image by detecting objects and other regions of interest, then aggregate features across these regions to produce a **pooled representation** richly expressing the image semantics. This feature vector is taken as input by a hierarchical recurrent neural network composed of two levels: a sentence RNN and a word RNN. The sentence RNN receives the image features, decides how many sentences to generate in the resulting paragraph, and produces an input topic vector for each sentence. Given this topic vector, the word RNN generates the words of a single sentence. We also show how to transfer knowledge from a dense image captioning [11] task to our model for paragraph generation.

### 4.1. Region Detector

The region detector receives an input image of size  $3 \times H \times W$ , detects regions of interest, and produces a feature vector of dimension  $D = 4096$  for each region. Our region detector follows [26, 11]; we provide a summary here for completeness: The image is resized so that its longest edge is 720 pixels, and is then passed through a convolutional network initialized from the 16-layer VGG network [28]. The resulting feature map is processed by a **region proposal network** [26], which regresses from a set of anchors to pro-

<sup>1</sup>Available at <http://cs.stanford.edu/people/ranjaykrishna/im2p/index.html>



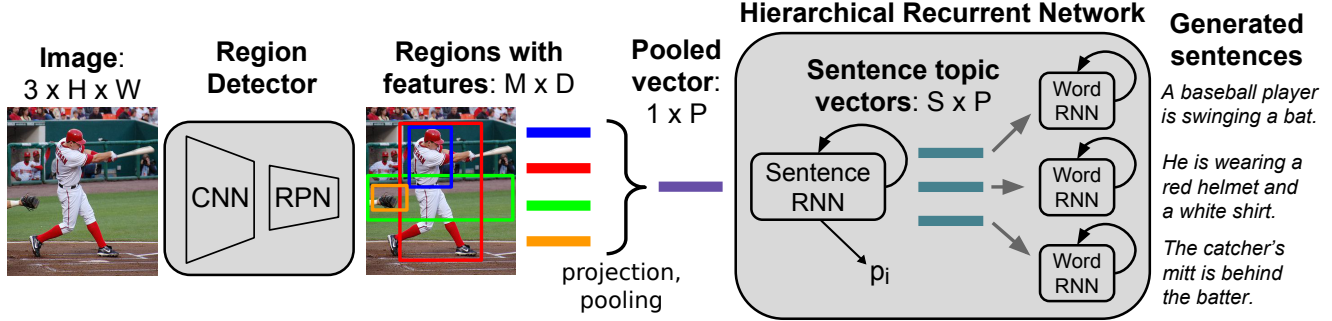


Figure 2. Overview of our model. Given an image (left), a region detector (comprising a convolutional network and a region proposal network) detects regions of interest and produces features for each. Region features are projected to  $\mathbb{R}^P$ , pooled to give a compact image representation, and passed to a hierarchical recurrent neural network language model comprising a sentence RNN and a word RNN. The sentence RNN determines the number of sentences to generate based on the halting distribution  $p_i$  and also generates sentence topic vectors, which are consumed by each word RNN to generate sentences.

pose regions of interest. These regions are projected onto the convolutional feature map, and the corresponding region of the feature map is reshaped to a fixed size using **bilinear interpolation** and processed by **two fully-connected layers** to give a vector of dimension  $D$  for each region.

Given a dataset of images and ground-truth regions of interest, the region detector can be trained in an end-to-end fashion as in [26] for object detection and [11] for dense captioning. Since paragraph descriptions do not have annotated groundings to regions of interest, we use a region detector trained for dense image captioning on the Visual Genome dataset [16], using the publicly available implementation of [11]. This produces  $M = 50$  detected regions.

One alternative worth noting is to use a region detector trained strictly for object detection, rather than dense captioning. Although such an approach would capture many salient objects in an image, its paragraphs would suffer: **an ideal paragraph describes not only objects, but also scenery and relationships**, which are better captured by dense captioning task that captures *all* noteworthy elements of a scene.

#### 4.2. Region Pooling

The region detector produces a set of vectors  $v_1, \dots, v_M \in \mathbb{R}^D$ , each describing a different region in the input image. We wish to aggregate these vectors into a single pooled vector  $v_p \in \mathbb{R}^P$  that compactly describes the content of the image. To this end, we **learn a projection matrix**  $W_{pool} \in \mathbb{R}^{P \times D}$  and bias  $b_{pool} \in \mathbb{R}^P$ ; the pooled vector  $v_p$  is computed by projecting each region vector using  $W_{pool}$  and taking an **elementwise maximum**, so that  $v_p = \max_{i=1}^M (W_{pool}v_i + b_{pool})$ . While alternative approaches for representing collections of regions, such as spatial attention [31], may also be possible, we view these as complementary to the model proposed in this paper; furthermore we note recent work [25] which has proven max pooling sufficient for representing any continuous set function, giving motivation that max pooling does not, in principle, sacrifice expressive power.

#### 4.3. Hierarchical Recurrent Network

The pooled region vector  $v_p \in \mathbb{R}^P$  is given as input to a **hierarchical neural language model** composed of two modules: a *sentence RNN* and a *word RNN*. The sentence RNN is responsible for deciding the number of sentences  $S$  that should be in the generated paragraph and for producing a  $P$ -dimensional *topic vector* for each of these sentences. Given a topic vector for a sentence, the word RNN generates the words of that sentence. We adopt the **standard LSTM** architecture [9] for both the word RNN and sentence RNN.

As an alternative to this hierarchical approach, one could instead use a non-hierarchical language model to directly generate the words of a paragraph, treating the end-of-sentence token as another word in the vocabulary. Our hierarchical model is advantageous because it reduces the length of time over which the recurrent networks must reason. Our paragraphs contain an average of **67.5 words** (Tab. 1), so a non-hierarchical approach must reason over dozens of time steps, which is extremely difficult for language models. However, since our paragraphs contain an average of **5.7 sentences**, each with an average of **11.9 words**, both the paragraph and sentence RNNs need only reason over much shorter time-scales, making learning an appropriate representation much more tractable.

**Sentence RNN** The sentence RNN is a **single-layer LSTM** with hidden size  $H = 512$  and **initial hidden and cell states set to zero**. At each time step, the sentence RNN receives the pooled region vector  $v_p$  as input, and in turn produces a sequence of hidden states  $h_1, \dots, h_S \in \mathbb{R}^H$ , one for each sentence in the paragraph. Each hidden state  $h_i$  is used in two ways: First, a linear projection from  $h_i$  and a logistic classifier produce a distribution  $p_i$  over the two states  $\{\text{CONTINUE} = 0, \text{STOP} = 1\}$  which determine whether the  $i$ th sentence is the last sentence in the paragraph. Second, the hidden state  $h_i$  is fed through a **two-layer fully-connected network to produce the topic vector  $t_i \in \mathbb{R}^P$  for the  $i$ th sentence of the paragraph, which is the input to the word RNN.**

**Word RNN** The word RNN is a **two-layer LSTM** with hidden size  $H = 512$ , which, given a topic vector  $t_i \in \mathbb{R}^P$  from the sentence RNN, is responsible for generating the words of a sentence. We follow the **input formulation** of [30]: the first and second inputs to the RNN are the topic vector and a special START token, and subsequent inputs are learned embedding vectors for the words of the sentence. At each timestep the hidden state of the last LSTM layer is used to predict a distribution over the words in the vocabulary, and a special END token signals the end of a sentence. After each Word RNN has generated the words of their respective sentences, these sentences are finally concatenated to form the generated paragraph.

#### 4.4. Training and Sampling

Training data consists of pairs  $(x, y)$ , with  $x$  an image and  $y$  a ground-truth paragraph description for that image, where  $y$  has  $S$  sentences, the  $i$ th sentence has  $N_i$  words, and  $y_{ij}$  is the  $j$ th word of the  $i$ th sentence. After computing the pooled region vector  $v_p$  for the image, we unroll the sentence RNN for  $S$  timesteps, giving a distribution  $p_i$  over the  $\{\text{CONTINUE}, \text{STOP}\}$  states for each sentence. We feed the sentence topic vectors to  $S$  copies of the word RNN, unrolling the  $i$ th copy for  $N_i$  timesteps, producing distributions  $p_{ij}$  over each word of each sentence. Our training loss  $\ell(x, y)$  for the example  $(x, y)$  is a weighted sum of two **cross-entropy** terms: a *sentence loss*  $\ell_{\text{sent}}$  on the stopping distribution  $p_i$ , and a *word loss*  $\ell_{\text{word}}$  on the word distribution  $p_{ij}$ :

$$\ell(x, y) = \lambda_{\text{sent}} \sum_{i=1}^S \ell_{\text{sent}}(p_i, \mathbf{I}[i = S]) \quad (1)$$

$$+ \lambda_{\text{word}} \sum_{i=1}^S \sum_{j=1}^{N_i} \ell_{\text{word}}(p_{ij}, y_{ij}) \quad (2)$$

To generate a paragraph for an image, we run the sentence RNN forward until the stopping probability  $p_i(\text{STOP})$  exceeds a threshold  $T_{\text{STOP}}$  or after  $S_{\text{MAX}}$  sentences, whichever comes first. We then sample sentences from the word RNN, choosing the most likely word at each timestep and stopping after choosing the STOP token or after  $N_{\text{MAX}}$  words. We set the parameters  $T_{\text{STOP}} = 0.5$ ,  $S_{\text{MAX}} = 6$ , and  $N_{\text{MAX}} = 50$  based on validation set performance.

#### 4.5. Transfer Learning

Transfer learning has become pervasive in computer vision. For tasks such as object detection [26] and image captioning [6, 12, 30, 31], it has become standard practice not only to process images with convolutional neural networks, but also to initialize the weights of these networks from weights that had been tuned for image classification, such as the 16-layer VGG network [28]. Initializing from a pre-trained convolutional network allows a form of knowledge

transfer from large classification datasets, and is particularly effective on datasets of limited size. Might transfer learning also be useful for paragraph generation?

We propose to utilize transfer learning in two ways. First, we initialize our region detection network from a model trained for dense image captioning [11]; although our model is end-to-end differentiable, we **keep this sub-network fixed** during training both for efficiency and also to prevent over-fitting. Second, we initialize the word embedding vectors, recurrent network weights, and output linear projection of the word RNN from a language model that had been trained on region-level captions [11], fine-tuning these parameters during training to be better suited for the task of paragraph generation. Parameters for tokens not present in the region model are initialized from the parameters for the UNK token. This initialization strategy allows our model to utilize linguistic knowledge learned on large-scale region caption datasets [16] to produce better paragraph descriptions, and we validate the efficacy of this strategy in our experiments.

### 5. Experiments

In this section we describe our paragraph generation experiments on the collected data described in Sec. 3, which we divide into 14,575 training, 2,487 validation, and 2,489 testing images.

#### 5.1. Baselines

**Sentence-Concat:** To demonstrate the difference between sentence-level and paragraph captions, this baseline samples and concatenates five sentence captions from a model [12] trained on MS COCO captions [20]. The first sentence uses **beam search** (beam size = 2) and the rest are **sampled**. The motivation for this is as follows: the image captioning model first produces the sentence that best describes the image as a whole, and subsequent sentences use sampling in order to generate a diverse range of sentences, since the alternative is to repeat the same sentence from beam search. We have validated that this approach works better than using either only beam search or only sampling, as the intent is to make the strongest possible comparison at a task-level to standard image captioning. We also note that, while Sentence-Concat is trained on **MS COCO**, all images in our dataset are also in MS COCO, and our descriptions were also written by users on **Amazon Mechanical Turk**.

**Image-Flat:** This model uses a **flat representation** for both images and language, and is equivalent to the standard image captioning model NeuralTalk [12]. It takes the whole image as input, and decodes into a paragraph token by token. We use the publically available implementation of [12], which uses the 16-layer VGG network [28] to extract CNN features and projects them as input into an LSTM [9], training the whole model jointly end-to-end.

	METEOR	CIDEr	BLEU-1	BLEU-2	BLEU-3	BLEU-4
Sentence-Concat	12.05	6.82	31.11	15.10	7.56	3.98
Template	14.31	12.15	37.47	21.02	12.30	7.38
DenseCap-Concat	12.66	12.51	33.18	16.92	8.54	4.54
Image-Flat ([12])	12.82	11.06	34.04	19.95	12.20	7.71
Regions-Flat-Scratch	13.54	11.14	37.30	21.70	13.07	8.07
Regions-Flat-Pretrained	14.23	12.13	38.32	22.90	14.17	<b>8.97</b>
Regions-Hierarchical (ours)	<b>15.95</b>	<b>13.52</b>	<b>41.90</b>	<b>24.11</b>	<b>14.23</b>	8.69
Human	19.22	28.55	42.88	25.68	15.55	9.66

Table 2. Main results for generating paragraphs. Our Region-Hierarchical method is compared with six baseline models and human performance along six language metrics.

**Template:** This method represents a very different approach to generating paragraphs, similar in style to an open-world version of more classical methods like BabyTalk [17], which converts a structured representation of an image into text via a handful of manually specified templates. The first step of our template-based baseline is to detect and describe many regions in a given target image using a pre-trained dense captioning model [11], which produces a set of region descriptions tied with bounding boxes and detection scores. The region descriptions are parsed into a set of subjects, verbs, objects, and various modifiers according to part of speech tagging and a handful of TokensRegex [3] rules, which we find suffice to parse the vast majority ( $\geq 99\%$ ) of the fairly simplistic and short region-level descriptions.

Each parsed word is scored by the sum of its detection score and the log probability of the generated tokens in the original region description. Words are then merged into a coherent graph representing the scene, where each node combines all words with the same text and overlapping bounding boxes. Finally, text is generated using the top  $N = 25$  scored nodes, prioritizing subject-verb-object triples first in generation, and representing all other nodes with existential “there is/are” statements.

**DenseCap-Concat:** This baseline is similar to Sentence-Concat, but instead concatenates DenseCap [11] predictions as separate sentences in order to form a paragraph. The intent of analyzing this method is to disentangle two key parts of the Template method: captioning and detection (*i.e.* DenseCap), and heuristic recombination into paragraphs. We combine the top  $n = 14$  outputs of DenseCap to form DenseCap-Concat’s output based on validation CIDEr+METEOR.

**Other Baselines:** “Regions-Flat-Scratch” uses a flat language model for decoding and initializes it from scratch. The language model input is the projected and pooled region-level image features. “Regions-Flat-Pretrained” uses a pre-trained language model. These baselines are included to show the benefits of decomposing the image into regions and pre-training the language model.

## 5.2. Implementation Details

All baseline neural language models use **two layers** of LSTM [9] units with 512 dimensions. The feature pooling dimension  $P$  is 1024, and we set  $\lambda_{sent} = 5.0$  and  $\lambda_{word} = 1.0$  based on validation set performance. Training is done via stochastic gradient descent with **Adam** [14], implemented in Torch. Of critical note is that model checkpoint selection is based on the best combined METEOR and CIDEr score on the validation set – although models tend to **minimize validation loss** fairly quickly, it takes much longer training for METEOR and CIDEr scores to stop improving.

## 5.3. Main Results

We present our main results at generating paragraphs in Tab. 2, which are evaluated across six language metrics: CIDEr [29], METEOR [5], and BLEU- $\{1,2,3,4\}$  [24]. The Sentence-Concat method performs poorly, achieving the lowest scores across all metrics. Its lackluster performance provides further evidence of the stark differences between single-sentence captioning and paragraph generation. Surprisingly, the hard-coded template-based approach performs reasonably well, particularly on CIDEr, METEOR, and BLEU-1, where it is competitive with some of the neural approaches. This makes sense: the template approach is provided with a strong prior about image content since it receives region-level captions [11] as input, and the many expletive “there is/are” statements it makes, though uninteresting, are safe, resulting in decent scores. However, its relatively poor performance on BLEU-3 and BLEU-4 highlights the limitation of reasoning about regions in isolation – it is unable to produce much text relating regions to one another, and further suffers from a lack of “connective tissue” that transforms paragraphs from a series of disconnected thoughts into a coherent whole. DenseCap-Concat scores worse than Template on all metrics except CIDEr, illustrating the necessity of Template’s caption parsing and recombination.

Image-Flat, trained on the task of paragraph generation, outperforms Sentence-Concat, and the region-based reasoning of Regions-Flat-Scratch improves results further on all metrics. Pre-training results in improvements on all met-

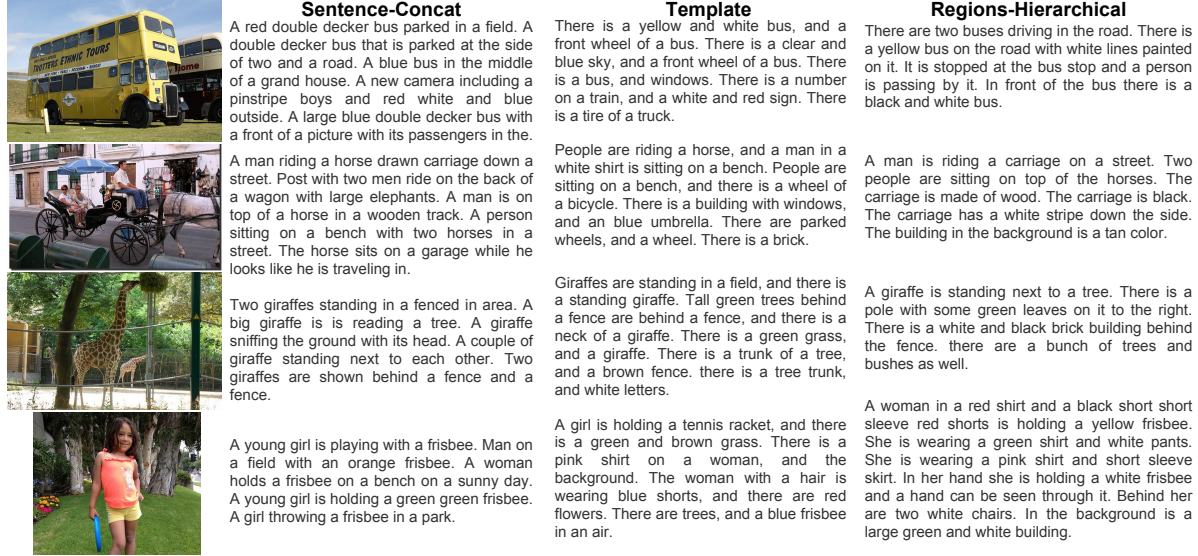


Figure 3. Example paragraph generation results for our model (Regions-Hierarchical) and the Sentence-Concat and Template baselines. The first three rows are positive results and the last row is a failure case.

rics, and our full model, Regions-Hierarchical, achieves the highest scores among all methods on every metric except BLEU-4. One hypothesis for the mild superiority of Regions-Flat-Pretrained on BLEU-4 is that it is better able to reproduce words immediately at the end and beginning of sentences more exactly due to their non-hierarchical structure, providing a slight boost in BLEU scores.

To make these metrics more interpretable, we performed a human evaluation by collecting an additional paragraph for 500 randomly chosen images, with results in the last row of Tab. 2. As expected, humans produce superior descriptions to any automatic method, performing better on all language metrics considered. Of particular note is the large gap between humans and the best model on CIDEr and METEOR, which are both designed to correlate well with human judgment [29, 5].

Finally, we note that we have also tried the SPICE evaluation metric [1], which has shown to correlate well with human judgements for sentence-level image captioning. Unfortunately, SPICE does not seem well-suited for evaluating long paragraph descriptions – it does not handle coreference or distinguish between different instances of the same object category. These are reasonable design decisions for sentence-level captioning, but is less applicable to paragraphs. In fact, human paragraphs achieved a considerably lower SPICE score than automated methods.

## 5.4. Qualitative Results

We present qualitative results from our model and the Sentence-Concat and Template baselines in Fig. 3. Some interesting properties of our model’s predictions include its use of coreference in the first example (“a bus”, “it”,

“the bus”) and its ability to capture relationships between objects in the second example. Also of note is the order in which our model chooses to describe the image: the first sentence tends to be fairly high level, middle sentences give some details about scene elements mentioned earlier in the description, and the last sentence often describes something in the background, which other methods are not able to capture. Anecdotally, we observed that this follows the same order with which most humans tended to describe images.

The failure case in the last row highlights another interesting phenomenon: even though our model was wrong about the semantics of the image, calling the girl “a woman”, it has learned that “woman” is consistently associated with female pronouns (“she”, “she”, “her hand”, “behind her”).

It is also worth noting the general behavior of the two baselines. Paragraphs from Sentence-Concat tend to be repetitive in sentence structure and are often simply inaccurate due to the sampling required to generate multiple sentences. On the other hand, the Template baseline is largely accurate, but has uninteresting language and lacks the ability to determine which things are most important to describe. In contrast, Regions-Hierarchical stays relevant and furthermore exhibits more interesting patterns of language.

5.5. Paragraph Language Analysis

To shed a quantitative light on the linguistic phenomena generated, in Tab. 3 we show statistics of the language produced by a representative spread of methods.

Our hierarchical approach generates text of similar average length and variance as human descriptions, with Sentence-Concat and the Template approach somewhat shorter and less varied in length. Sentence-Concat is also

There are two buses driving in the road. There is a yellow bus on the road with white lines painted on it. It is stopped at the bus stop and a person is passing by it. In front of the bus there is a black and white bus.

A man is riding a carriage on a street. Two people are sitting on top of the horses. The carriage is made of wood. The carriage is black. The carriage has a white stripe down the side. The building in the background is a tan color.

A giraffe is standing next to a tree. There is a pole with some green leaves on it to the right. There is a white and black brick building behind the fence. There are a bunch of trees and bushes as well.

A woman in a red shirt and a black short short sleeve red shorts is holding a yellow frisbee. She is wearing a green shirt and white pants. She is wearing a pink shirt and short sleeve skirt. In her hand she is holding a white frisbee and a hand can be seen through it. Behind her are two white chairs. In the background is a large green and white building.



	Average Length	Std. Dev. Length	Diversity	Nouns	Verbs	Pronouns	Vocab Size
Sentence-Concat	56.18	4.74	34.23	32.53	9.74	0.95	2993
Template	60.81	7.01	45.42	23.23	11.83	0.00	422
Regions-Hierarchical	70.47	17.67	40.95	24.77	13.53	2.13	1989
Human	67.51	25.95	69.92	25.91	14.57	2.42	4137

Table 3. Language statistics of test set predictions. Part of speech statistics are given as percentages, and diversity is calculated as in Section 3. “Vocab Size” indicates the number of unique tokens output across the entire test set, and human numbers are calculated from ground truth. Note that the diversity score for humans differs slightly from the score in Tab. 1, which is calculated on the entire dataset.

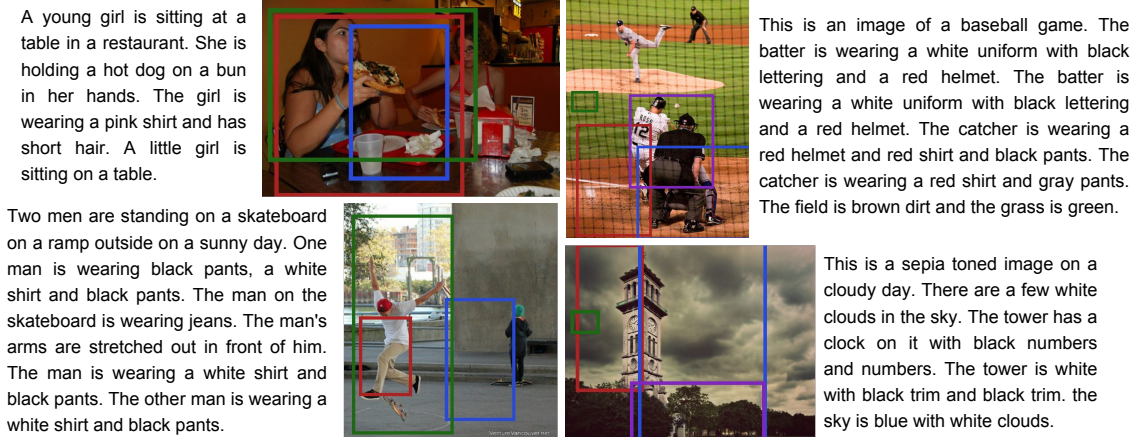


Figure 4. Examples of paragraph generation from only a few regions. Since only a small number of regions are used, this data is extremely out of sample for the model, but it is still able to focus on the regions of interest while ignoring the rest of the image.

the least diverse method, though all automatic methods remain far less diverse than human sentences, indicating ample opportunity for improvement. According to this diversity metric, the Template approach is actually the most diverse automatic method, which may be attributed to how the method is hard-coded to sequentially describe each region in the scene in turn, regardless of importance or how interesting such an output may be (see Fig. 3). While both our hierarchical approach and the Template method produce text with similar portions of nouns and verbs as human paragraphs, only our approach was able to generate a reasonable quantity of pronouns. Our hierarchical method also had a much wider vocabulary compared to the Template approach, though Sentence-Concat, trained on hundreds of thousands of MS COCO [20] captions, is a bit larger.

## 5.6. Generating Paragraphs from Fewer Regions

As an exploratory experiment in order to highlight the interpretability of our model, we investigate generating paragraphs from a smaller number of regions than the  $M = 50$  used in the rest of this work. Instead, we only give our method access to the top few detected regions as input, with the hope that the generated paragraph focuses only on those particularly regions, preferring not to describe other parts of the image. The results for a handful of images are shown in Fig. 4. Although the input is extremely out of sample com-

pared to the training data, the results are still quite reasonable – the model generates paragraphs describing the detected regions without much mention of objects or scenery outside of the detections. Taking the top-right image as an example, despite a few linguistic mistakes, the paragraph generated by our model mentions the batter, catcher, dirt, and grass, which all appear in the top detected regions, but does not pay heed to the pitcher or the umpire in the background.

## 6. Conclusion

In this paper we have introduced the task of describing images with long, descriptive paragraphs, and presented a hierarchical approach for generation that leverages the compositional structure of both images and language. We have shown that paragraph generation is different from traditional image captioning and have tailored our model to suit these differences. Experimentally, we have demonstrated the advantages of our approach over traditional image captioning methods and shown how region-level knowledge can be effectively transferred to paragraph captioning. We have also demonstrated the benefits of our model in interpretability, generating descriptive paragraphs using only a subset of image regions. We anticipate further opportunities for knowledge transfer at the intersection of vision and language, and project that visual and lingual compositionality will continue to lie at the heart of effective paragraph generation.



## References

- [1] P. Anderson, B. Fernando, M. Johnson, and S. Gould. Spice: Semantic propositional image caption evaluation. In *European Conference on Computer Vision*, pages 382–398. Springer, 2016.
- [2] W. Chan, N. Jaitly, Q. V. Le, and O. Vinyals. Listen, attend, and spell: A neural network for large vocabulary conversational speech recognition. In *ICASSP*, 2016.
- [3] A. X. Chang and C. D. Manning. Tokensregex: Defining cascaded regular expressions over tokens. Technical report, CSTR 2014-02, Department of Computer Science, Stanford University, 2014.
- [4] X. Chen and C. Lawrence Zitnick. Mind’s eye: A recurrent visual representation for image caption generation. In *CVPR*, 2015.
- [5] M. Denkowski and A. Lavie. Meteor universal: Language specific translation evaluation for any target language. In *EACL Workshop on Statistical Machine Translation*, 2014.
- [6] J. Donahue, L. Anne Hendricks, S. Guadarrama, M. Rohrbach, S. Venugopalan, K. Saenko, and T. Darrell. Long-term recurrent convolutional networks for visual recognition and description. In *CVPR*, 2015.
- [7] S. El Hihi and Y. Bengio. Hierarchical recurrent neural networks for long-term dependencies. In *NIPS*, 1995.
- [8] A. Farhadi, M. Hejrati, M. A. Sadeghi, P. Young, C. Rashtchian, J. Hockenmaier, and D. Forsyth. Every picture tells a story: Generating sentences from images. In *ECCV*, 2010.
- [9] S. Hochreiter and J. Schmidhuber. Long short-term memory. *Neural computation*, 1997.
- [10] M. Hodosh, P. Young, and J. Hockenmaier. Framing image description as a ranking task: Data, models and evaluation metrics. *Journal of Artificial Intelligence Research*, 47:853–899, 2013.
- [11] J. Johnson, A. Karpathy, and L. Fei-Fei. DenseCap: Fully convolutional localization networks for dense captioning. In *CVPR*, 2016.
- [12] A. Karpathy and L. Fei-Fei. Deep visual-semantic alignments for generating image descriptions. In *CVPR*, 2015.
- [13] A. Karpathy, A. Joulin, and L. Fei-Fei. Deep fragment embeddings for bidirectional image sentence mapping. In *NIPS*, 2014.
- [14] D. Kingma and J. Ba. Adam: A method for stochastic optimization. In *ICLR*, 2015.
- [15] J. Koutnik, K. Greff, F. Gomez, and J. Schmidhuber. A clockwork RNN. In *ICML*, 2014.
- [16] R. Krishna, Y. Zhu, O. Groth, J. Johnson, K. Hata, J. Kravitz, S. Chen, Y. Kalantidis, L.-J. Li, D. A. Shamma, M. Bernstein, and L. Fei-Fei. Visual genome: Connecting language and vision using crowdsourced dense image annotations. *arXiv preprint arXiv:1602.07332*, 2016.
- [17] G. Kulkarni, V. Premraj, S. Dhar, S. Li, Y. Choi, A. C. Berg, and T. L. Berg. Baby talk: Understanding and generating image descriptions. In *CVPR*, 2011.
- [18] J. Li, M.-T. Luong, and D. Jurafsky. A hierarchical neural autoencoder for paragraphs and documents. In *ACL*, 2015.
- [19] R. Lin, S. Liu, M. Yang, M. Li, M. Zhou, and S. Li. Hierarchical recurrent neural network for document modeling. In *EMNLP*, 2015.
- [20] T.-Y. Lin, M. Maire, S. Belongie, J. Hays, P. Perona, D. Ramanan, P. Dollár, and C. L. Zitnick. Microsoft coco: Common objects in context. In *ECCV*, 2014.
- [21] C. D. Manning, M. Surdeanu, J. Bauer, J. R. Finkel, S. Bethard, and D. McClosky. The stanford CoreNLP natural language processing toolkit. In *ACL (System Demonstrations)*, pages 55–60, 2014.
- [22] J. Mao, W. Xu, Y. Yang, J. Wang, Z. Huang, and A. Yuille. Deep captioning with multimodal recurrent neural networks (m-RNN). *ICLR*, 2015.
- [23] M. P. Marcus, M. A. Marcinkiewicz, and B. Santorini. Building a large annotated corpus of english: The penn treebank. *Computational linguistics*, 19(2):313–330, 1993.
- [24] K. Papineni, S. Roukos, T. Ward, and W.-J. Zhu. Bleu: a method for automatic evaluation of machine translation. In *ACL*, pages 311–318. Association for Computational Linguistics, 2002.
- [25] C. R. Qi, H. Su, K. Mo, and L. J. Guibas. Pointnet: Deep learning on point sets for 3d classification and segmentation. *arXiv preprint arXiv:1612.00593*, 2016.
- [26] S. Ren, K. He, R. Girshick, and J. Sun. Faster r-cnn: Towards real-time object detection with region proposal networks. In *NIPS*, 2015.
- [27] A. Rohrbach, M. Rohrbach, W. Qiu, A. Friedrich, M. Pinkal, and B. Schiele. Coherent multi-sentence video description with variable level of detail. In *German Conference on Pattern Recognition*, pages 184–195. Springer, 2014.
- [28] K. Simonyan and A. Zisserman. Very deep convolutional networks for large-scale image recognition. In *ICLR*, 2015.
- [29] R. Vedantam, C. Lawrence Zitnick, and D. Parikh. CIDEr: Consensus-based image description evaluation. In *CVPR*, 2015.
- [30] O. Vinyals, A. Toshev, S. Bengio, and D. Erhan. Show and tell: A neural image caption generator. In *CVPR*, 2015.
- [31] K. Xu, J. Ba, R. Kiros, K. Cho, A. Courville, R. Salakhudinov, and Y. Bengio. Show, attend, and tell: Neural image caption generation with visual attention. In *ICML*, 2015.
- [32] L. Yao, A. Torabi, K. Cho, N. Ballas, C. Pal, H. Larochelle, and A. Courville. Describing videos by exploiting temporal structure. In *ICCV*, 2015.
- [33] Q. You, H. Jin, Z. Wang, C. Fang, and J. Luo. Image captioning with semantic attention. In *CVPR*, 2016.
- [34] P. Young, A. Lai, M. Hodosh, and J. Hockenmaier. From image descriptions to visual denotations: New similarity metrics for semantic inference over event descriptions. *Transactions of the Association for Computational Linguistics*, 2:67–78, 2014.
- [35] H. Yu, J. Wang, Z. Huang, Y. Yang, and W. Xu. Video paragraph captioning using hierarchical recurrent neural networks. In *CVPR*, 2016.